

# MPI – Message passing interface

Advanced Programming in C/C++


14 . 01 . 2019

BY: Narges Mehran & Dr. Dragi Kimovski




# Shared vs. Distributed Memory Programming

Shared Memory – All tasks access the same memory, hence the same data - **OpenMP**.



Distributed Memory – All memory is local. Data sharing is by explicitly transporting data from one task to another – **MPI**.



# MPI

- Message Passing Interface (MPI) is a message-passing standard created to function on a wide variety of parallel architectures.
- It is standardized in collaboration between industry and academia.
- There are multiple implementations of the standard, such as:
  - MPICH: <http://www.mpich.org/>
  - OpenMPI: <https://www.open-mpi.org/>

# A Simple Example

- A serial program

```
#include<stdio.h>
```

```
#define PID 0
```

```
main()
```

```
{
```

```
    int i;
```

```
    printf("Greetings from process %d!\n", PID);
```

```
}
```

Output:      Greetings from process 0



# A Simple Example (cont.)

- A parallel program using MPI

```
#include<mpi.h>
```

```
main(int argc, char** argv){
```

```
...
```

```
    MPI_Init(&argc, &argv); /* Initialize MPI */
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &size); /* Get the number of  
processes */
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); /* Get my process number  
(rank) */
```

Parallel Region

```
    MPI_Finalize(); /* Exit MPI */
```

```
}
```

# A Simple Example (cont.)

- `MPI_Comm_size`: Determines the size of the group associated with a communicator (the number of processes)
  - `int MPI_Comm_size(MPI_Comm MPI_COMM_WORLD, int *size)`
- `MPI_Comm_rank`: Determines the rank of the calling process in the communicator (a number between 0 and size-1)
  - `int MPI_Comm_rank(MPI_Comm MPI_COMM_WORLD, int *rank)`
- Sending a message:
  - `int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm MPI_COMM_WORLD)`
  - “buf” is the address of the message to be sent, with “count” elements of type “datatype”
- Receiving a message:
  - `int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm MPI_COMM_WORLD, MPI_Status *status)`

# A Simple Example (cont.)

- A parallel program using MPI (cont.)

```
....  
int rank, size, tag, rc, i;  
MPI_Status status;  
char message[20];  
if (rank==0)  
{  
    strcpy(message, "Hello, world");  
    for (i=1;i<size;++i)  
        rc = MPI_Send(message, 13, MPI_CHAR, i, tag,  
MPI_COMM_WORLD);  
}  
else  
    rc = MPI_Recv(message, 13, MPI_CHAR, 0, tag, MPI_COMM_WORLD,  
&status);  
printf("Greetings from process %d of %d\n", rank, size);
```



# A Simple Example (cont.)

- A parallel program using MPI (cont.)

*Output:*

Greetings from process 0 of 3

Greetings from process 1 of 3

Greetings from process 2 of 3



# MPICH

- A freely available, portable implementation of MPI on Linux clusters
- Primarily developed by Argonne National Laboratory in the USA
- <http://www.mpich.org/>
- MPI Routines and Constants:
  - <https://www.mpich.org/static/docs/v3.3/>

# Installation

- Installing MPICH:

```
>>> tar -xzf mpich-3.3.tar.gz
```

```
>>> cd mpich-3.3
```

```
>>> ./configure --disable-fortran
```

```
>>> make; sudo make install
```

```
>>> mpiexec --version
```

HYDRA build details: Version: 3.3

Release Date: Wed Nov 21 11:32:40 CST 2018

CC: gcc

CXX: g++

F77:

F90:

# Compilation

- GNU Compiler Collection (GCC) and MPICH
- In Linux Terminal, enter one of the following commands:
  - ❑ `mpicc mpi_code.c -o mpi_code.out`
  - ❑ `mpicxx mpi_code.cc -o mpi_code.out`



# Compilation (cont.)

```
#include <stdio.h>
#include "mpi.h"
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    printf("Hello world \n");
    MPI_Finalize();
}
```

```
program Hello_World
include 'mpif.h'
integer ierror
call MPI_INIT(ierror)
print *, 'Hello World'
call MPI_FINALIZE(ierror)
end
```

```
#include "mpi.h"
#include <iostream>
int main(int argc, char** argv) {
    MPI::Init(argc, argv);
    cout <<"Hello world \n";
    MPI::Finalize();
}
```

## Compilers

C

\$ mpicc

Fortran

\$ mpif90

C++

\$ mpicxx

# Execution

- MPI (using multi-core processors)
  - `mpdboot`
  - `mpiexec -n #processes ./ mpi_code.out`
  - `mpirun`
- Running hellow
  - `mpiexec -n 3 --hosts Narges,localhost hellow.out`
  - `mpirun -np 3 --hosts Narges,localhost ./hellow.out`

# Let's start coding your program!


1. Write an MPI based program which starts communicating a message (like a number) back and forth from process 0 to process 1
  - Such as an SMS between your friend and you
    - ❖ A point-to-point communication
  - Use `MPI_Send` and `MPI_Recv` methods



# A parallel program using MPI\_BCAST

- One processor sends some data to all processors in a group
- The message is sent from the root process to all processes in the group, including the root process
  - `int MPI_Bcast(void *buf, int count, MPI_Datatype datatype, int root, MPI_Comm MPI_COMM_WORLD)`
- To synchronize all processes within a communicator
- A node calling it will be blocked until all nodes within the group have called it:
  - `int MPI_Barrier(MPI_Comm MPI_COMM_WORLD)`

# A parallel program using MPI\_BCAST (cont.)



<http://mpitutorial.com>

# A parallel program using MPI\_BCAST (cont.)

...

```
MPI_Bcast(data, num_elements, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Barrier(MPI_COMM_WORLD);
total_my_bcast_time -= MPI_Wtime();
MPI_Barrier(MPI_COMM_WORLD);
total_mpi_bcast_time += MPI_Wtime();
if ( my_rank == 0)
    // Print off timing information if (my_rank == 0)
    {
        printf("Data size = %d \n", num_elements * (int)sizeof(int) );
        printf("Avg MPI_Bcast time = %f\n", total_mpi_bcast_time);
    }
```



# Use MPI\_BCAST to write a new program!

- **MPI\_Bcast** distributes data from one process (the root) to all others in a communicator
2. Extend your previous program for three processes and use broadcasting
- It can be found in this link:  
<http://mpitutorial.com/tutorials/mpi-broadcast-and-collective-communication/>

# References

- 1) Gropp, William D., et al. Using MPI: portable parallel programming with the message-passing interface. Vol. 1. MIT press, 1999.
- 2) Parallelization with OpenMP and MPI A Simple Example (C)
- 3) <https://www.mpi-forum.org/>
- 4) <https://www.mpich.org/documentation/guides/>
- 5) Parallel programming - with MPI and OpenMP
- 6) <http://mpitutorial.com/tutorials/>



# Happy parallel programming.