

Docker Containers Tutorial- part3


Docker Swarm

Distributed Systems UE,

Docker swarmkit

- [Swarmkit](#): Cluster management and orchestration features in Docker Engine 1.12 or later.
- It turns a pool of Docker hosts into a single, virtual host.
- When Swarmkit is enabled, we call Docker Engine running in swarm mode.
- See the feature list: [Swarm mode overview](#).
- This project focuses on micro-service architecture.
- It supports service reconciliation, load balancing, service discovery, built-in certificate rotation, etc.

Clustering



<https://docs.docker.com/>

Docker Nodes?

Machines running *SwarmKit* can be grouped together in order to form a *Swarm*, coordinating tasks with each other. Once a machine joins, it becomes a *Swarm Node*. Nodes can either be *worker* nodes or *manager* nodes.

- **Worker Nodes** are responsible for running Tasks using an *Executor*. *SwarmKit* comes with a default *Docker Container Executor* that can be easily swapped out.
- **Manager Nodes** on the other hand accept specifications from the user and are responsible for reconciling the desired state with the actual cluster state.

Features of SwarmKit

Some of *SwarmKit*'s main features are:

- Orchestration
- Scheduling
- Cluster management

Orchestration

- **Desired State Reconciliation:** *SwarmKit* constantly compares the desired state against the current cluster state and reconciles the two if necessary. For instance, if a node fails, *SwarmKit* reschedules its tasks onto a different node.
- **Service Types:** There are different types of services. The project currently ships with two of them out of the box
 - ✓ **Replicated Services** are scaled to the desired number of replicas.
 - ✓ **Global Services** run one task on every available node in the cluster.

Orchestration (cont.)

- **Configurable Updates:** At any time, you can change the value of one or more fields for a service. After you make the update, *SwarmKit* reconciles the desired state by ensuring all tasks are using the desired settings. By default, it performs a lockstep update - that is, update all tasks at the same time. This can be configured through different knobs:
 - ✓ **Parallelism** defines how many updates can be performed at the same time.
 - ✓ **Delay** sets the minimum delay between updates. *SwarmKit* will start by shutting down the previous task, bring up a new one, wait for it to transition to the *RUNNING* state *then* wait for the additional configured delay. Finally, it will move onto other tasks.
- **Restart Policies:** The orchestration layer monitors tasks and reacts to failures based on the specified policy. The operator can define restart conditions, delays and limits (maximum number of attempts in a given time window). *SwarmKit* can decide to restart a task on a different machine. This means that faulty nodes will gradually be drained of their tasks.

Scheduling

node attribute	matches	example
node.id	node's ID	node.id == 2ivku8v2gvtg4
node.hostname	node's hostname	node.hostname != node-2
node.ip	node's IP address	node.ip != 172.19.17.0/24
node.role	node's manager or worker role	node.role == manager
node.platform.os	node's operating system	node.platform.os == linux
node.platform.arch	node's architecture	node.platform.arch == x86_64
node.labels	node's labels added by cluster admins	node.labels.security == high
engine.labels	Docker Engine's labels	engine.labels.operatingsystem == ubuntu 14.04

Cluster Management

- **State Store:** Manager nodes maintain a strongly consistent, replicated (Raft based) and extremely fast (in-memory reads) view of the cluster which allows them to make quick scheduling decisions while tolerating failures.
- **Topology Management:** Node roles (*Worker / Manager*) can be dynamically changed through API/CLI calls.
- **Node Management:** An operator can alter the desired availability of a node: Setting it to *Paused* will prevent any further tasks from being scheduled to it while *Drained* will have the same effect while also re-scheduling its tasks somewhere else (mostly for maintenance scenarios).

Docker Swarm

Setup using the hosted discovery service

- Create a cluster:
 - `$ swarm create`
- Add nodes to a cluster:
 - `$ swarm join --add=<node_ip>`
- Start Swarm
 - `$ swarm manage --addr=<swarm_ip>`

Scheduling

Scheduling in Docker Swarm relies on filters and strategies

- spread
- random
- binpack

A manager node, schedules the tasks on a set of worker nodes.

A survey of Docker Swarm scheduling strategies

Anwar Hassen at Aalto University

Scheduling (cont.)

- Consider that there is a client that gets connected to the server farm.
- Consider the following scenario in which the manager of the server farm receives requests from the client as follows:
 - It is supposed to schedule the received tasks on the worker nodes. The tasks have different execution times. Manager receives a list consisting of execution times of tasks that should be sent to the worker nodes.
 - The execution times are 16, 7, 20, 10, 11, 22, and 12 minutes. Each task is allowed to have the resources of worker node for 25 minutes.
 - Manager schedules the tasks by Bin-Packing Algorithm. This algorithm packs the tasks in two methods: first-fit and best-fit. What is the difference between two methods? Please implement both methods.

Scheduling (cont.)

- <https://gist.githubusercontent.com/vfarcic/750fc4117bad9d8619004081af171896/raw/3134886e92c09f47ac37edfc151d745fa8e4235e/02-docker-swarm.sh>

References

- <https://www.docker.com/>
- Bernstein, D., “Containers and cloud: From lxc to docker to kubernetes,” *IEEE Cloud Computing*, 1(3), pp.81-84, 2014.
- Scott McCarty. 2018. A Practical Introduction to Container Terminology. <https://developers.redhat.com/blog/2018/02/22/container-terminology-practical-introduction/>
- <http://docker.atbaker.me/index.html>
- Boza, E. F., Abad, C. L., Narayanan, S. P., Balasubramanian, B., & Jang, M. (2019, December). A Case for Performance-Aware Deployment of Containers. In *Proceedings of the 5th International Workshop on Container Technologies and Container Clouds* (pp. 25-30).
- https://dockerbook.com/TheDockerBook_sample.pdf